

PROFESSIONAL SUMMARY

Staff MLOps Engineer with 7 years specializing in LLM infrastructure at AI-native companies. Deep expertise in Amazon Bedrock, evaluation frameworks, and multi-cloud LLM deployment pipelines. Seeking to relocate; no current NY connections.

EXPERIENCE

Staff MLOps Engineer

Cohere | San Francisco, CA | 2020 – Present

- Leads MLOps infrastructure for enterprise customers deploying Cohere models on AWS and Azure; manages 30+ production LLM environments.
- Designed multi-environment CI/CD pipeline (dev/staging/prod) with automated evaluation gates measuring hallucination, coherence, and task completion.
- Built prompt versioning system and automated rollback pipelines; reduced production incidents by 60% in 18 months.
- Works daily with Amazon Bedrock, LangGraph, and custom evaluation frameworks for agentic workflows.
- Implements secrets management, IAM policies, and sandboxed agent environments per enterprise security requirements.

MLOps Engineer

Weights & Biases | San Francisco, CA | 2018 – 2020

- Built CI/CD integrations enabling automated experiment tracking and model evaluation at 500+ enterprise customers.
 - Developed automated model evaluation pipelines; contributed to open-source MLOps tooling.
-

EDUCATION

B.S. Computer Science, UC Berkeley, 2017

SKILLS

AWS (Bedrock, SageMaker, ECS, Lambda, CloudWatch, CodePipeline), Azure OpenAI, LangGraph, LangChain, Python, Terraform, Docker, CI/CD, evaluation frameworks, MLflow, GitHub Actions, secrets management