

PROFESSIONAL SUMMARY

MLOps Engineer specializing in LLM evaluation frameworks and production AI infrastructure. 4 years at JPMorgan Chase building hallucination measurement pipelines and AWS Bedrock deployment systems. Passionate about AI reliability engineering and agent observability.

EXPERIENCE

MLOps Engineer

JPMorgan Chase | New York, NY | 2021 – Present

- Led production deployment of LLM-based customer service agents using Amazon Bedrock; manages model lifecycle including versioning, evaluation gates, and rollback.
- Built comprehensive evaluation pipeline measuring hallucination rates, response coherence, and workflow reliability across 6 production AI systems.
- Designed automated rollback mechanisms triggered by SLO breaches; reduced mean time to recovery from 45 minutes to under 8 minutes.
- Implemented LangGraph multi-agent workflows for financial analysis tools; established governance guardrails limiting agent permissions.
- Owns CloudWatch and Grafana observability stack; defined 23 SLOs covering latency, error rate, and cost-per-request targets.

Junior ML Engineer

Capital One | New York, NY | 2019 – 2021

- Built AWS Glue and Step Functions pipelines for credit risk model feature engineering.
 - Supported SageMaker-based model deployment infrastructure for fraud detection models.
-

EDUCATION

M.S. Computer Science, New York University (Courant), 2019; B.Tech Computer Engineering, IIT Delhi, 2017

SKILLS

AWS (Bedrock, SageMaker, Lambda, Step Functions, CloudWatch, Glue), Python, LangGraph, LangChain, MLflow, CI/CD, Docker, Terraform, evaluation frameworks, Grafana, SLO design, incident response